

# CS 412: Spring '24

## Introduction To Data Mining

### Assignment 3

(Due Monday, April 8, 23:59)

- The homework is due on Monday, April 8, 2024, at 23:59. We are using Gradescope for all homework assignments. In case you haven't already, make sure to join this course on Gradescope using the code shared on Canvas. Contact the TAs if you face any technical difficulties while submitting the assignment. Please do NOT email a copy of your solution. We will NOT accept late submissions (without a reasonable justification).
- Please use Campuswire if you have questions about the homework. Make sure to appropriately tag your post. Also, scroll through previous posts to make sure that your query was not answered previously. In case you are sending us an email regarding this Assignment, start the subject with "CS 412 Spring '24 HW3:" and include **all** TAs and the Instructor (Jeffrey, Xinyu, Kowshika, Sayar, Ruby).
- Please write your code entirely by yourself. We will run MOSS to detect code similarity.
- All programming must be done in Python 3.
- The homework will be graded using Gradescope. You will be able to submit your code as many times as you want.
- The grade generated by the autograder upon submission will be your final grade for this assignment. There are no post deadline tests.
- Two python files named `homework3_q1.py`, `homework3_q2.py` containing starter code are available on Canvas.
- Do NOT add any additional `import` statements to your code. Required libraries have already been included in the starter code.
- For submitting on Gradescope, you would need to upload **both** the files:
  1. `homework3_q1.py`
  2. `homework3_q2.py`
- Late submission policy: there will be a 24-hour grace period without any grade reduction, i.e., Gradescope will accept late submissions until **Tuesday, April 9, 2024, at 23:59.** Unfortunately, we will NOT accept late submissions past the grace period (without a reasonable justification).

1. (35 points) This problem will focus on developing code for Bayes Theorem applied to estimating the type of the underlying candy bag based on candies drawn from the bag (similar to the example discussed in class).

There are five types of candy bags:

- $\pi_1$  fraction are  $h_1$ :  $p_1$  fraction “cherry” candies,  $(1 - p_1)$  fraction “lime” candies
- $\pi_2$  fraction are  $h_2$ :  $p_2$  fraction “cherry” candies,  $(1 - p_2)$  fraction “lime” candies
- $\pi_3$  fraction are  $h_3$ :  $p_3$  fraction “cherry” candies,  $(1 - p_3)$  fraction “lime” candies
- $\pi_4$  fraction are  $h_4$ :  $p_4$  fraction “cherry” candies,  $(1 - p_4)$  fraction “lime” candies
- $\pi_5$  fraction are  $h_5$ :  $p_5$  fraction “cherry” candies,  $(1 - p_5)$  fraction “lime” candies

For the specific example discussed in class:

$$\begin{aligned} \pi_1 = 0.1, \quad \pi_2 = 0.2, \quad \pi_3 = 0.4, \quad \pi_4 = 0.2, \quad \pi_5 = 0.1, \\ p_1 = 1, \quad p_2 = 0.75, \quad p_3 = 0.5, \quad p_4 = 0.25, \quad p_5 = 0. \end{aligned}$$

A bag is given to us. But we do not know which type of bag  $h \in \{h_1, h_2, h_3, h_4, h_5\}$  it is. We will be drawing a sequence of candies  $c_1, c_2, \dots, c_i \in \{\text{“cherry”}, \text{“lime”}\}$  from the given bag and, using Bayes rule, maintain posterior probabilities of the type of bag conditioned on the candies which have been drawn, i.e.,

$$\begin{aligned} \text{After drawing } c_1 : \quad & p(\pi_h | c_1), \quad h = 1, \dots, 5 \\ \text{After drawing } c_1, c_2 : \quad & p(\pi_h | c_1, c_2), \quad h = 1, \dots, 5 \\ \dots & \dots \end{aligned}$$

We will make the following assumptions regarding the setup:

- The probabilities  $p_h, h = 1, \dots, k$  of drawing “cherry” candies from the bag do not change as candies are being drawn the bag. As a result, the probabilities  $(1 - p_h), h = 1, \dots, k$  of drawing “lime” candies from the bag also do not change as candies are being drawn the bag.
- The joint probability of drawing different candies from a bag are conditionally independent given the bag, i.e.,

$$\begin{aligned} p(c_1, c_2 | \pi_h) &= p(c_1 | \pi_h) p(c_2 | \pi_h), \\ p(c_1, c_2, c_3 | \pi_h) &= p(c_1 | \pi_h) p(c_2 | \pi_h) p(c_3 | \pi_h), \end{aligned}$$

and so on.

You will have to develop code for the following function:

`my_Bayes_candy( $\pi, p, c_{1:10}$ )`, which uses prior probabilities  $\pi$ , conditional probabilities  $p$ , and sequence of 10 candy draws  $c_{1:10}$  to compute posterior probabilities of each type of bag.

The function will have the following **input**:

- Prior probability of each type of bag:  $\pi = [\pi_1, \pi_2, \pi_3, \pi_4, \pi_5]$
- Conditional probability of cherry candies in each type of bag:  $p = [p_1, p_2, p_3, p_4, p_5]$
- Sequence of 10 candies drawn from the bag under consideration:  $c_{1:10} = [c_1, c_2, \dots, c_{10}]$

where  $c_i \in \{0, 1\}$  with “0” denoting cherry and “1” denoting lime candy. For the specific example discussed in class:

$$c = [1, 1, 1, 1, 1, 1, 1, 1, 1, 1].$$

The function will have the following **output**:

- (a) A list of posterior probabilities of each type of bag  $\pi_h, h = 1, \dots, 5$  after every subsequence of candy draws, i.e.,  $c_{1:t} = [c_1, \dots, c_t], t \in \{1, \dots, 10\}$ .

In particular, the output will be the following posterior probabilities:

$$\begin{aligned} \text{After drawing } c_1 : & \quad p(\pi_h | c_1) , \quad h = 1, \dots, 5 \\ \text{After drawing } c_1, c_2 : & \quad p(\pi_h | c_1, c_2) , \quad h = 1, \dots, 5 \\ \dots & \quad \dots \\ \text{After drawing } c_1, \dots, c_{10} : & \quad p(\pi_h | c_1, \dots, c_{10}) , \quad h = 1, \dots, 5 \end{aligned}$$

The function `my_Bayes_candy` will return a two-dimensional list of size 10x5 containing the aforementioned posterior probabilities. Specifically, the returned list must contain probabilities in the format:

$$\begin{aligned} & [[p(\pi_1 | c_1), p(\pi_2 | c_1), \dots, p(\pi_5 | c_1)], \\ & \quad [p(\pi_1 | c_1, c_2), \dots, p(\pi_5 | c_1, c_2)], \\ & \quad \dots, \\ & \quad [p(\pi_1 | c_1, \dots, c_{10}), \dots, p(\pi_5 | c_1, \dots, c_{10})]] \end{aligned}$$

Note that your code should work for any given  $\pi, p, c_{1:10}$ . We will consider 5 types of bags and 2 types of candies in each bag. `my_Bayes_candy` would have a time limit of 10 seconds. You could use the specific example discussed in class to test out your implementation. However, your code must work for any valid input.

2. (65 points) The problem will focus on developing your own code for: ***K*-fold cross validation**.

Your code will be evaluated using seven standard classification models applied to a multi-class classification dataset.

**Dataset:** We will be using the following dataset for the assignment.

**Digits:** The Digits dataset comes prepackaged with **scikit-learn** (`sklearn.datasets.load_digits`). The dataset has 1797 points, 64 features, and 10 classes corresponding to ten numbers  $0, 1, \dots, 9$ . The dataset was (likely) created from the following dataset:

<https://archive.ics.uci.edu/ml/datasets/Optical+Recognition+of+Handwritten+Digits>

**Classification Methods.** We will consider seven classification methods from **scikit-learn** and **xgboost**:

- Gaussian Naive Bayes Classifier: **GaussianNB**,
- Linear Support Vector Classifier: **LinearSVC**,
- Support Vector Classifier: **SVC**,
- Decision Tree Classifier: **DecisionTreeClassifier**,
- AdaBoost Classifier: **AdaBoostClassifier**,
- Gradient Boosting Classifier: **XGBClassifier**, and
- Feed-forward Neural Network: **MLPClassifier**.

Use the following parameters for these methods (do not specify any additional parameters):

- **GaussianNB**: No parameters to be specified
- **LinearSVC**: `dual='auto', max_iter=2048, random_state=412`
- **SVC**: `gamma='scale', C=10, random_state=412`
- **DecisionTreeClassifier**: `max_depth=5, random_state=412`
- **AdaBoostClassifier**: `algorithm='SAMME', random_state=412`
- **XGBClassifier**: `max_depth=6, random_state=412`
- **MLPClassifier**: `alpha=1, max_iter=64, shuffle=False, random_state=412`

You will have to develop **code** for the following two functions:

- (a) **get\_splits(*n*, *k*, *seed*)**, which returns: randomized *k* ‘almost equal’ sized lists of unique disjoint indices from the set of all indices  $\{0, \dots, n - 1\}$ , where the randomization depends on the integer **seed**. By ‘almost equal’, we mean the cardinality of the lists can differ by at most 1. These *k* list of indices correspond to the *k* folds over which cross-validation will be performed. The function will have the following **output**:

- i. a python list containing exactly *k* lists. Each of these sublists should be disjoint, each of size roughly  $\frac{n}{k}$ , contain elements from the set  $\{0, 1, \dots, n - 1\}$  and must not contain repeated elements. The union of all the *k* sublists should include all elements in  $\{0, \dots, n - 1\}$ .

Input **seed** determines the randomization and the output should be the same every time we use the same **seed** for a given *n*, *k*, and should be (randomly) different for different values of the **seed**.

For example, **get\_splits**(4, 2, 1) may return `[[0,2], [1,3]]`, and the output must be same every time with the same input; **get\_splits**(4, 2, 73) may return `[[0,3], [1,2]]`; **get\_splits**(7, 2, 2) may return `[[0,2,4,6], [1,3,5]]`; **get\_splits**(11, 3, 7) may return `[[0,3,6,9], [1,4,7,10], [2,5,8]]`.

For our tests, *n* would be less than 1200. **get\_splits** would have a time limit of 10 seconds.

- (b) **my\_cross\_val(method, X, y, splits)**, which runs *k*-fold cross-validation for **method** on the dataset (*X*, *y*). The **input parameters** are:

- i. **method**: a python string which specifies the (class) name of one of the seven classification methods under consideration,
- ii. **X, y**: the data for the classification problem
- iii. **splits**: the output of the `get_splits` method (`len(splits) = k`)

The function will have the following **output**:

- i. the model's **accuracy** for each of the  $k$  folds.

The (auto)grader will judge your solution as correct if the difference between the reported and the expected mean **accuracy** is within  $10^{-3}$ . `my_cross_val` would have a time limit of 2 minutes.

Within `my_cross_val`, strictly use the *splits* encoded in the input parameter **splits**. Do not define your splits for  $K$ -fold cross validation within this method.

**Important:** Make sure that you are NOT inadvertently shuffling your training data during  $K$ -fold cross validation. The training examples within any particular split should be in the same order as the input  $X$ .

Use `my_cross_val` to return the error rates in each fold for  $k$  fold cross-validation for the specific **method** (one of the seven methods specified above) with the parameters outlined above.

To verify your code locally or on Google Colab, you could use:

```
from sklearn.datasets import load_digits
digits = load_digits()
X, y = digits.data, digits.target
```